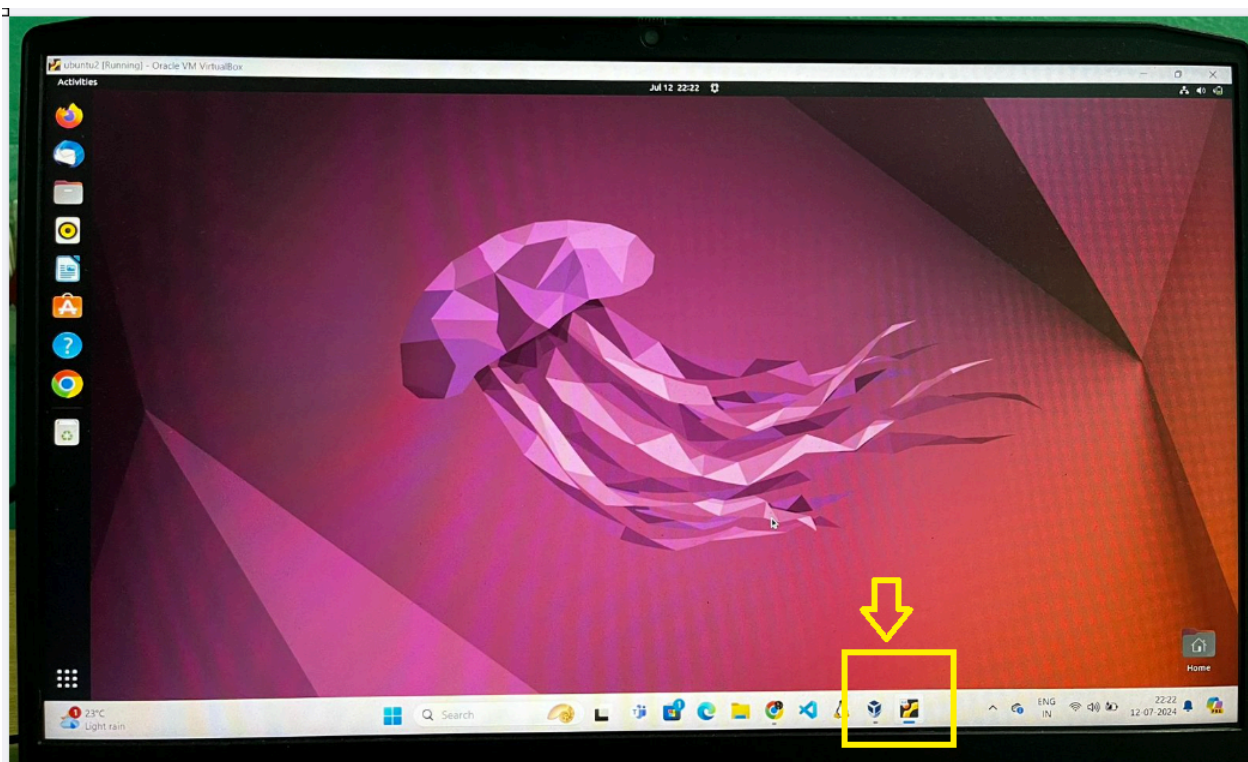
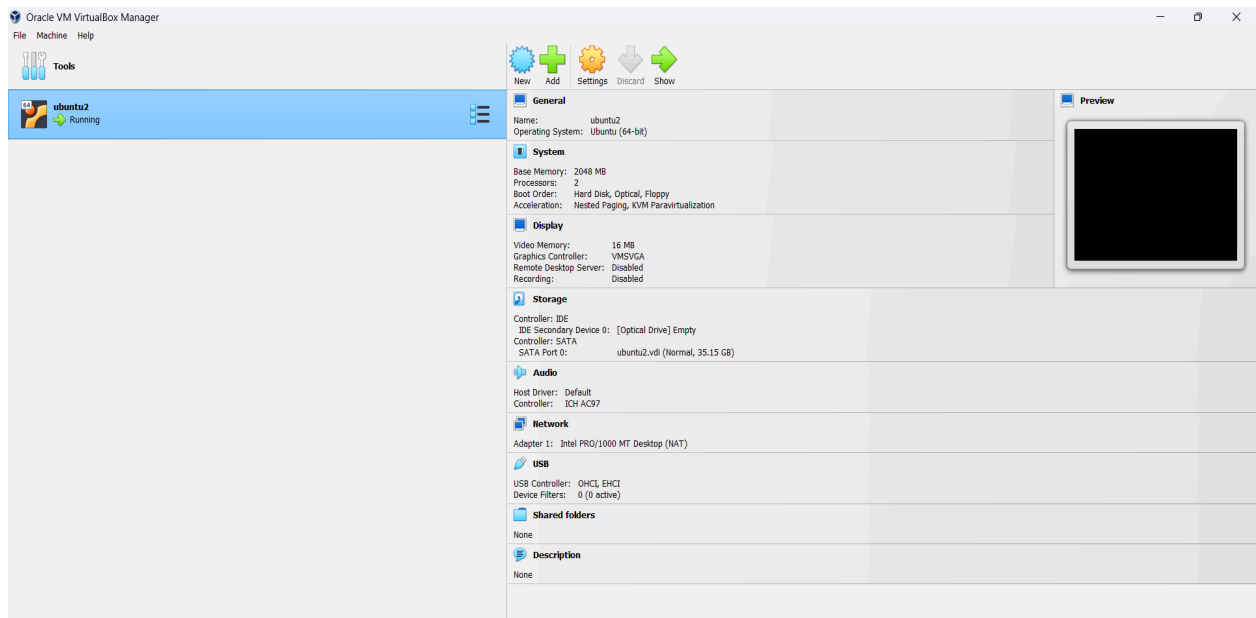


## WEEK 10 - WEEK 12 ASSIGNMENT

Submitted By : Karthik S Deshpande

### 1. Host a Ubuntu Virtual Machine using Oracle VM Virtual Box

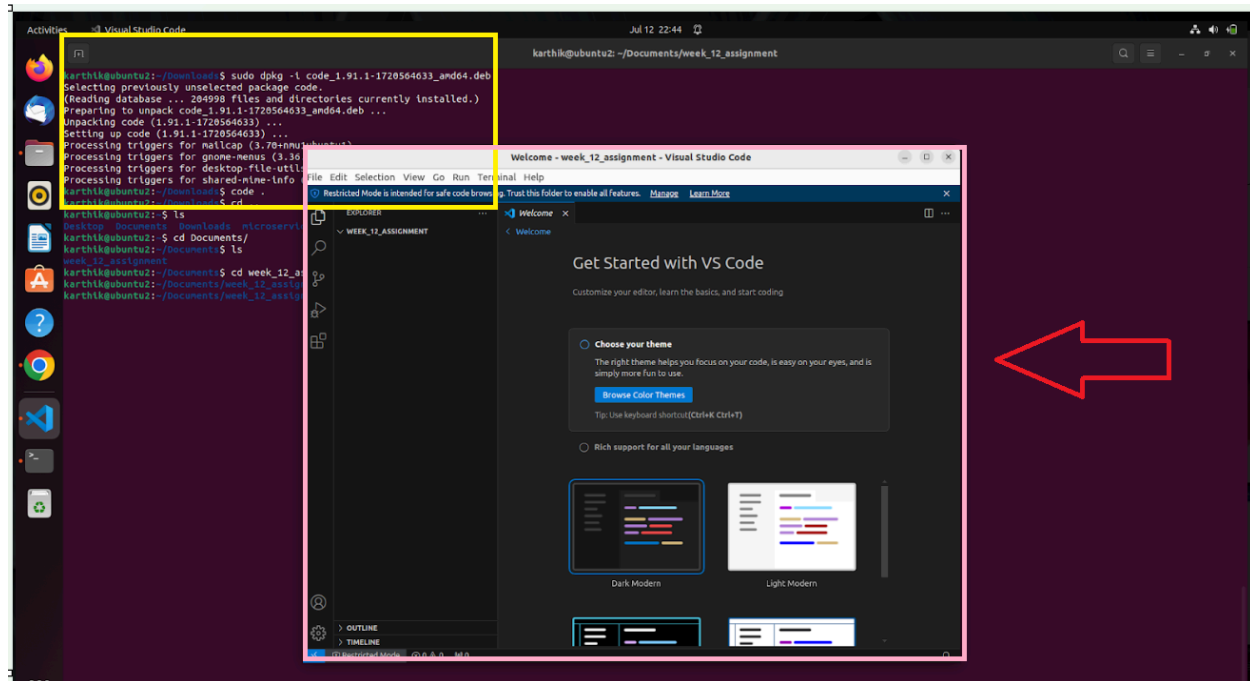
Solution : Hosted Ubuntu Virtual machine using oracle VM Virtual box



## 2. Set up Visual Studio code on Ubuntu VM.

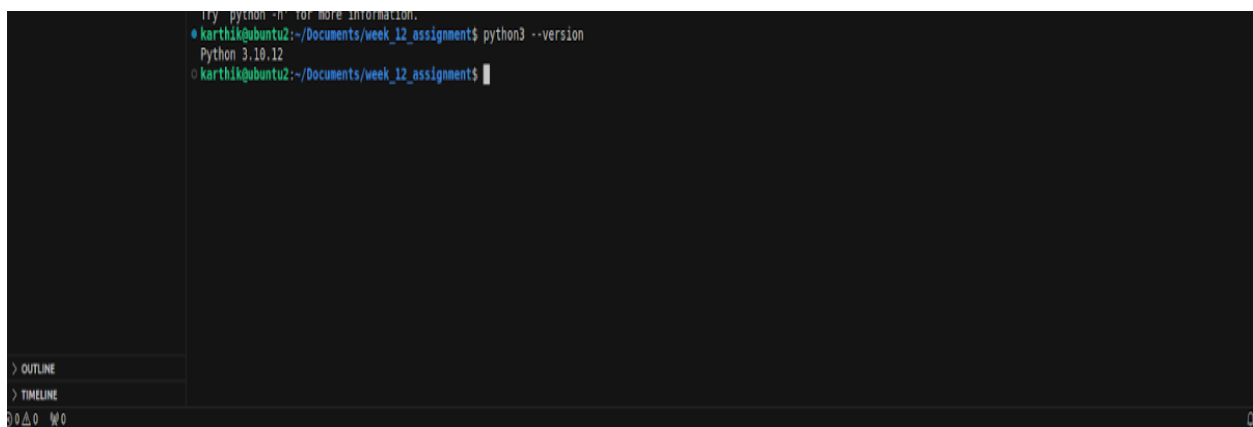
Solution : Downloaded the VSCode installation file from <https://code.visualstudio.com>

In the terminal of Linux ran command to install VSCODE.



## 3. Set up Python

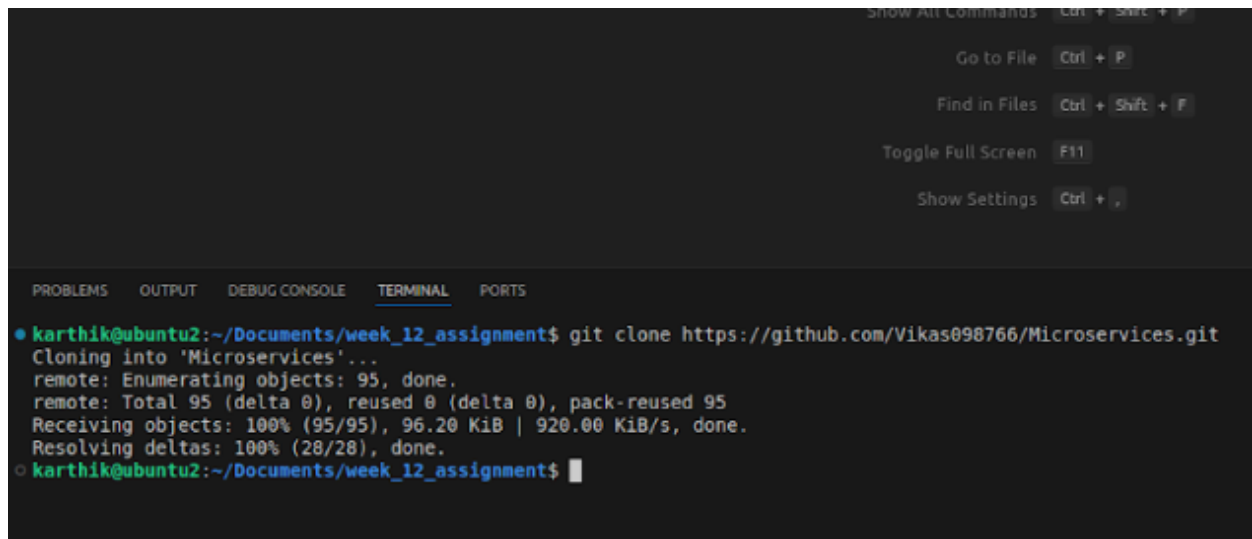
Solution : Python is set up, please find below snap with python version



#### 4. Clone this Github repository <https://github.com/Vikas098766/Microservices.git>

**Solution :** Cloned using the command

- `git clone https://github.com/Vikas098766/Microservices.git`



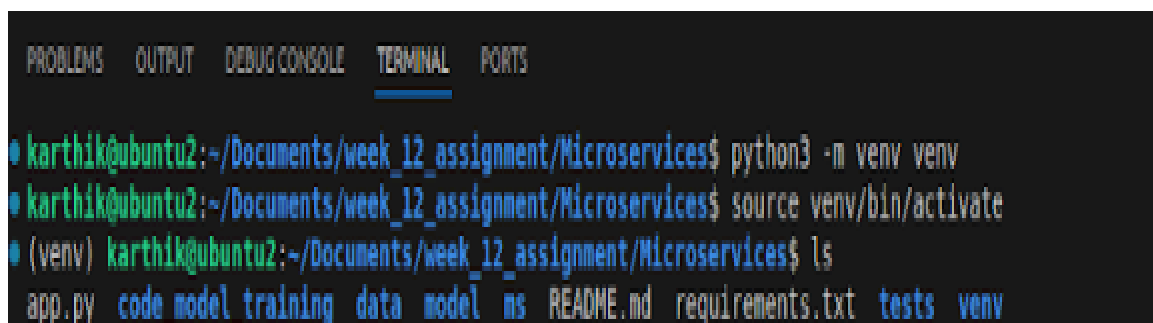
The screenshot shows a terminal window with the following content:

```
karthik@ubuntu2:~/Documents/week_12_assignment$ git clone https://github.com/Vikas098766/Microservices.git
Cloning into 'Microservices'...
remote: Enumerating objects: 95, done.
remote: Total 95 (delta 0), reused 0 (delta 0), pack-reused 95
Receiving objects: 100% (95/95), 96.20 KiB | 920.00 KiB/s, done.
Resolving deltas: 100% (28/28), done.
karthik@ubuntu2:~/Documents/week_12_assignment$
```

#### 5. Create a Virtual Environment.

**Solution :** Created Virtual Environment using commands

- `python3 -m venv venv`
- `source venv/bin/activate`



The screenshot shows a terminal window with the following content:

```
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ python3 -m venv venv
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ source venv/bin/activate
(venv) karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ ls
app.py  code_model_training  data  model  ns  README.md  requirements.txt  tests  venv
```

## 6. Install the dependencies from requirements.txt file.

**Solution :** Installed all dependencies present in requirements.txt file using the command

- Pip install -r requirements.txt

```
(venv) karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ pip install -r requirements.txt
Collecting click==8.0.3
  Using cached click-8.0.3-py3-none-any.whl (97 kB)
Collecting cycler==0.11.0
  Using cached cycler-0.11.0-py3-none-any.whl (6.4 kB)
Collecting Flask==2.0.2
  Using cached Flask-2.0.2-py3-none-any.whl (95 kB)
Collecting fonttools==4.28.5
  Using cached fonttools-4.28.5-py3-none-any.whl (890 kB)
Collecting gunicorn==20.1.0
  Using cached gunicorn-20.1.0-py3-none-any.whl (79 kB)
Collecting itsdangerous==2.0.1
  Using cached itsdangerous-2.0.1-py3-none-any.whl (18 kB)
Collecting Jinja2==3.0.3
  Using cached Jinja2-3.0.3-py3-none-any.whl (133 kB)
Collecting joblib==1.1.0
  Using cached joblib-1.1.0-py2.py3-none-any.whl (306 kB)
Collecting kiwisolver==1.3.2
  Using cached kiwisolver-1.3.2-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (1.6 MB)
Collecting MarkupSafe==2.0.1
  Using cached MarkupSafe-2.0.1-cp310-cp310-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2010_x86_64.whl (30 kB)
Collecting matplotlib==3.5.1
  Using cached matplotlib-3.5.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.9 MB)
Collecting numpy==1.22.0
  Using cached numpy-1.22.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (16.8 MB)
Collecting packaging==21.3
  Using cached packaging-21.3-py3-none-any.whl (40 kB)
Collecting pandas==1.3.5
  Using cached pandas-1.3.5-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.5 MB)
Collecting Pillow==9.0.0
```

## 7. Train and save the model.

**Solution :** Trained and saved the model.

- Command : python code\_model\_training/train.py

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
(venv) karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ python code_model_training/train.py
Accuracy: 0.9736842105263158
<sklearn.metrics.plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7ca36440f310>
/home/karthik/Documents/week_12_assignment/Microservices/code_model_training/train.py:54: UserWarning: Matplotlib is currently using agg, which is a non-GUI backend, so cannot show the figure.
  plt.show()
```

## 8. Test the Flask web application.

**Solution :** Tested web application by running the command.

- flask run -p 5000

```
plt.show()
(venv) karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ flask run -p 5000
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)
```

## 9. Test the application and make predictions using the example calls available in the folder /tests

**Solution :** Tested the the end point /info

**Command :** curl -X GET <http://localhost:5000/info>

**Command :** curl -X GET <http://localhost:5000/health>

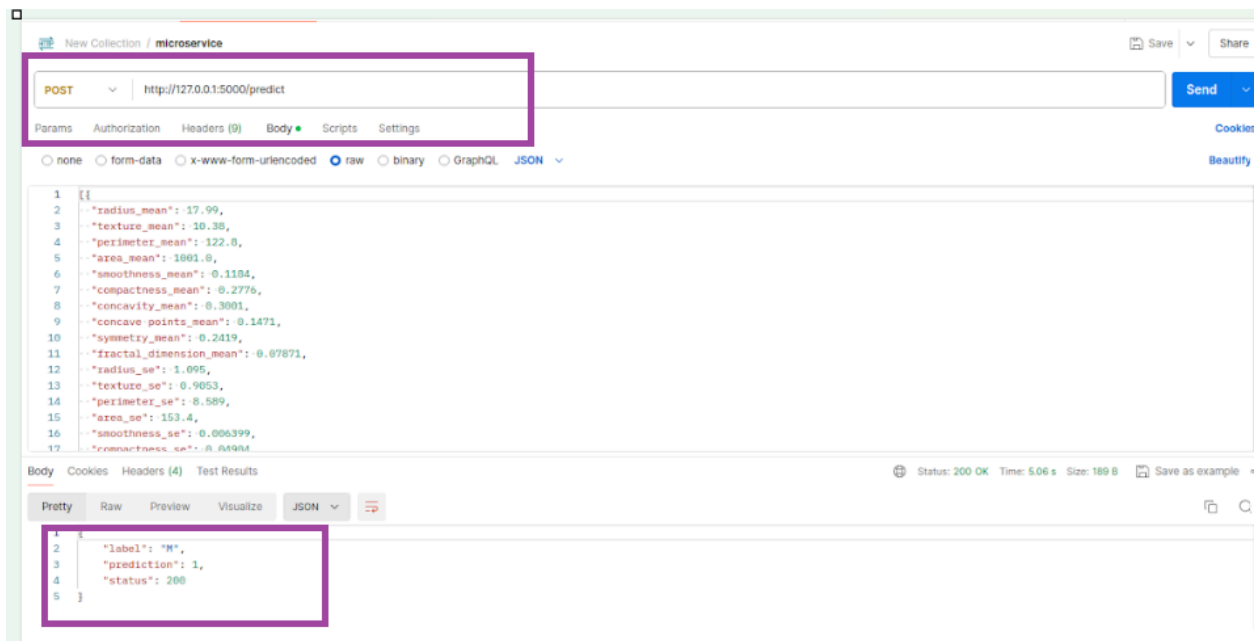
```
{ "label": "M", "prediction": "M", "status": "200" }
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ curl -X GET http://localhost:5000/info
{"name": "Breast Cancer Wisconsin (Diagnostic)", "version": "v1.0.0"}
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ curl -X GET http://localhost:5000/health
okkarthik@ubuntu2:~/Documents/week_12_assignment/Microservices$
```

**Command :** curl -d '{"radius\_mean": 17.99, "texture\_mean": 10.38, "perimeter\_mean": 122.8, "area\_mean": 1001.0, "smoothness\_mean": 0.1184, "compactness\_mean": 0.2776, "concavity\_mean": 0.3001, "concave points\_mean": 0.1471, "symmetry\_mean": 0.2419, "fractal\_dimension\_mean": 0.07871, "radius\_se": 1.095, "texture\_se": 0.9053, "perimeter\_se": 8.589, "area\_se": 153.4, "smoothness\_se": 0.006399, "compactness\_se": 0.04904, "concavity\_se": 0.05373, "concave points\_se": 0.01587, "symmetry\_se": 0.03003, "fractal\_dimension\_se": 0.006193, "radius\_worst": 25.38, "texture\_worst": 17.33, "perimeter\_worst": 184.6, "area\_worst": 2019.0, "smoothness\_worst": 0.1622, "compactness\_worst": 0.6656, "concavity\_worst": 0.7119, "concave points\_worst": 0.2654, "symmetry\_worst": 0.4601, "fractal\_dimension\_worst": 0.1189}]' \ -H "Content-Type: application/json" \ -X POST http://0.0.0.0:5000/predict

```
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ curl -d '{"radius_mean": 17.99, "texture_mean": 10.38, "perimeter_mean": 122.8, "area_mean": 1001.0, "smoothness_mean": 0.1104, "compactness_mean": 0.2776, "concavity_mean": 0.3001, "concave points_mean": 0.1471, "symmetry_mean": 0.2419, "fractal_dimension_mean": 0.07871, "radius_se": 1.095, "texture_se": 0.9053, "perimeter_se": 8.589, "area_se": 153.4, "smoothness_se": 0.006399, "compactness_se": 0.04904, "concavity_se": 0.05373, "concave points_se": 0.01507, "symmetry_se": 0.03003, "fractal_dimension_se": 0.006193, "radius_worst": 25.38, "texture_worst": 17.33, "perimeter_worst": 184.0, "area_worst": 2019.0, "smoothness_worst": 0.1022, "compactness_worst": 0.0650, "concavity_worst": 0.7119, "concave points_worst": 0.2654, "symmetry_worst": 0.4601, "fractal_dimension_worst": 0.1189}]' \
-H 'Content-Type: application/json' \
-X POST http://0.0.0.0:5000/predict
{"label": "M", "prediction": 1, "status": 200}
```

API ENDPOINT with /predict got the output as {"label": "M", "prediction": 1, "status": 200}

Verified using postman for more visibility



**10. Create a docker image containing everything needed to run the application.**


**Solution : Steps to create a docker image.**

1. Created the text file named dockerfile using the command as **touch dockerfile**

```
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ touch dockerfile
```



2. Within the txt file adding the following content within it



```
1 # Use an official Python runtime as a parent image
2 FROM python:3.9-slim
3
4 # Set the working directory inside the container
5 WORKDIR /usr/src/app
6
7 # Copy the requirements file into the container
8 COPY requirements.txt ./
9
10 # Install dependencies
11 RUN pip install --no-cache-dir -r requirements.txt
12
13 # Copy the rest of the application code into the container
14 COPY . .
15
16 # Expose the port the app runs on
17 EXPOSE 5000
18
19 # Define the command to run the app
20 CMD ["flask", "run", "--host=0.0.0.0", "--port=5000"]
```

3. Build the docker image with the name as **my-python-app**

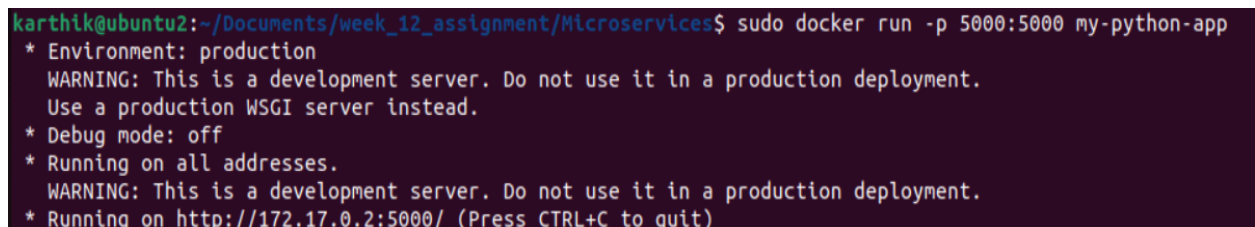
- Command : **sudo docker build -t my-python-app .**



```
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ sudo docker build -t my-python-app .
[+] Building 246.8s (10/10) FINISHED
=> [internal] load build definition from dockerfile
=> => transferring dockerfile: 538B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/python:3.9-slim
=> [1/5] FROM docker.io/library/python:3.9-slim@sha256:a6c12ec09f13df9d4b8b4e4d08678c1b212d89885be14b6c72b633bee2a520f4
=> => resolve docker.io/library/python:3.9-slim@sha256:a6c12ec09f13df9d4b8b4e4d08678c1b212d89885be14b6c72b633bee2a520f4
=> => sha256:a6c12ec09f13df9d4b8b4e4d08678c1b212d89885be14b6c72b633bee2a520f4 10.41kB / 10.41kB
=> => sha256:4719115deb9cc7a5479a7d3c57cfceac2be89fcaf0fed8c747e8dfb4b01a79a3 1.94kB / 1.94kB
=> => sha256:b97320a8c1caf64deeebb911ff8eb75bf12f671408a85302dd33b5ede2d1cd1 6.90kB / 6.90kB
=> => sha256:f11c1adaa26e078479ccdd45312ea3b88476441b91be0ec898a7e07bfd05badc 29.13MB / 29.13MB
=> => sha256:c1f67e58a3d2a9d9c5f38c8c3fc029ff3b9d6e0045b935c99e9ffc4182070fa1 3.51MB / 3.51MB
=> => sha256:9370038d11852cad5a70691e76b0ddc8e669018bc770cad15c23a3def629b874 11.89MB / 11.89MB
```

4. Run the Docker Container

- Command : **sudo docker run -p 5000:5000 my-python-app**



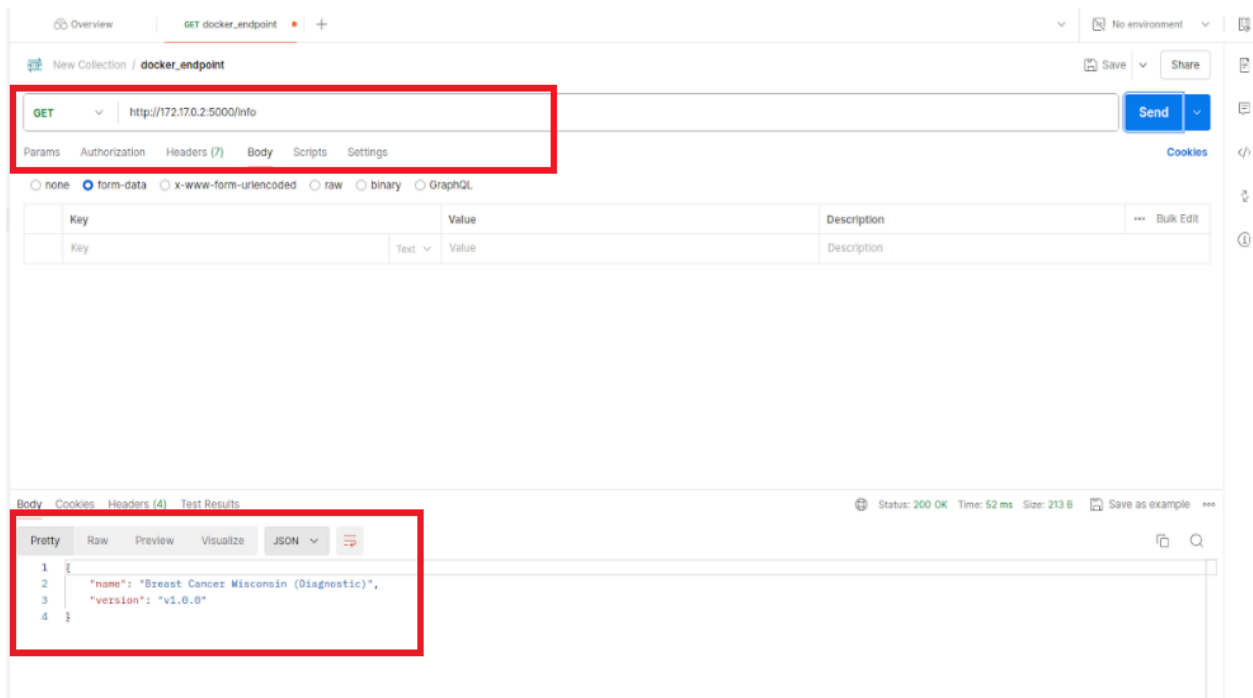
```
karthik@ubuntu2:~/Documents/week_12_assignment/Microservices$ sudo docker run -p 5000:5000 my-python-app
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on all addresses.
  WARNING: This is a development server. Do not use it in a production deployment.
* Running on http://172.17.0.2:5000/ (Press CTRL+C to quit)
```

Docker image is running successfully.

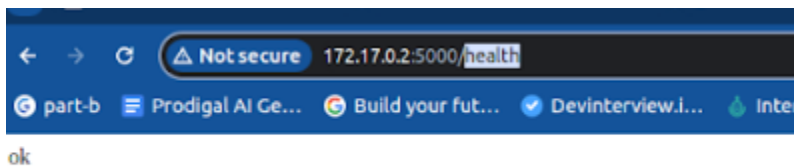
11. Run the containerized application as a prediction service and test it locally by passing some example calls and get the prediction.

**Solution :** To check the Docker image service locally with the help of POSTMAN end points as

- **/info**

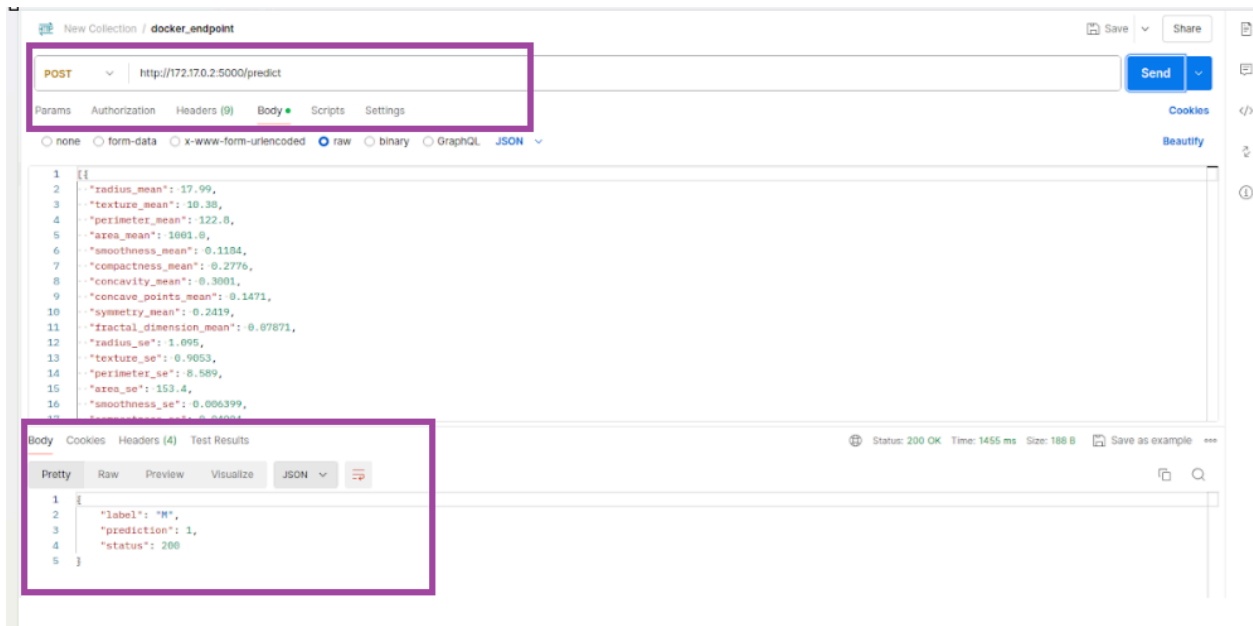


- **/health**





- /predict



Passed parameters as

```
[
  {
    "radius_mean": 17.99,
    "texture_mean": 10.38,
    "perimeter_mean": 122.8,
    "area_mean": 1001.0,
    "smoothness_mean": 0.1184,
    "compactness_mean": 0.2776,
    "concavity_mean": 0.3001,
    "concave_points_mean": 0.1471,
    "symmetry_mean": 0.2419,
    "fractal_dimension_mean": 0.07871,
    "radius_se": 1.095,
    "texture_se": 0.9053,
```

```
"perimeter_se": 8.589,  
"area_se": 153.4,  
"smoothness_se": 0.006399,  
"compactness_se": 0.04904,  
"concavity_se": 0.05373,  
"concave_points_se": 0.01587,  
"symmetry_se": 0.03003,  
"fractal_dimension_se": 0.006193,  
"radius_worst": 25.38,  
"texture_worst": 17.33,  
"perimeter_worst": 184.6,  
"area_worst": 2019.0,  
"smoothness_worst": 0.1622,  
"compactness_worst": 0.6656,  
"concavity_worst": 0.7119,  
"concave_points_worst": 0.2654,  
"symmetry_worst": 0.4601,  
"fractal_dimension_worst": 0.1189  
}  
]
```

\*\*\*\*\* THE END \*\*\*\*\*